

Simulation Experiments

Salvin Chwodhury

February 17, 2025

1 Introduction

The purpose of this paper is to experiment with the implications of sampling from a probability distribution. We start by importing, defining, evaluating, and plotting a single normal distribution. We then sample this distribution and analyze the relationships and trends related to sample size.

Additionally, we examine the effects of stochastic sampling, where we explore type I errors (false positives) and type II errors (false negatives). Finally, we use bootstrapping to estimate the confidence interval for the median of the dataset.

1.1 Dataset Description

For this paper, we generate two kinds of datasets. The first dataset is generated randomly by creating 2000 evenly spaced numbers between (mean - 5 times the standard deviation) and (mean + 5 times the standard deviation).

The second dataset is contained within a `dataset.csv` file. This `.csv` file consists of three variables: `variable_1`, `variable_2`, and `variable_3`. Each variable contains 500 observations, all of which are floating-point numbers that can be either positive or negative.

1.2 Dataset Summary Statistics

Referring to the `dataset.csv` file, we examined summary statistics by converting it into a `pandas` dataframe and using `.info()` and `.describe()`. Using both of these methods, we obtained the following summary statistics below.

x	v_1	v_2	v_3
Mean	-0.002	0.822	0.942
Median	0.032	0.797	0.938
Standard Deviation	0.996	1.006	1.004
Maximum	3.057	4.150	3.623
Minimum	-2.989	-2.428	-1.731
Lower Quartile	-0.666	0.176	0.303
Upper Quartile	0.649	1.498	1.551

Table 1: Statistical Summary of Variables

2 Methods

3 Results

3.1 Creating & Plotting a Probability Density Function

To begin, we start by creating the probability density function (PDF). We initialize a normal distribution object using `norm()` by passing in the desired mean and standard deviation values. Next, we use `.linspace()` to generate an array of equally spaced values ranging from `mean - 5*standard_deviation` to `mean + 5*standard_deviation`, producing 2000 points.

Following that, we proceed to plot the probability density function. Using the `matplotlib` library, we plot the X values (generated using `.linspace()`) against the corresponding density values, which are obtained by applying the `.pdf()` function to the normal distribution object.

What we ended up getting is the graph as shown below displayed as Figure 1. As we can see, there are a total 2000 values displayed and plotted as a Probability Density Function. The highest density occurs when the value of x is at 0 and the lowest densities occur when the the values of x are -4 or 4.




Figure 1: PDF of Normal Distribution Curve

3.2 Creating & Plotting Cumulative Distribution Function

Next, we move on to creating the cumulative distribution function. Using the normalized distribution object, which has a passed in mean and standard deviation parameters, and the array of equally spaced numbers using `.linspace()`, we then use `.cdf()` to calculate the cumulative distribution function.

We then move to plotting the cumulative distribution function. Using the `matplotlib` library, we then plot the X values that we have generated against the cumulative probabilities that were generated after applying the `.cdf` function.

What we ended up getting is the graph as shown below in Figure 2. As we can see, the values range from 0 to 1, and we end up getting a cumulative probability that goes from 0 and up to 1, hence producing the curve as shown. Similar to figure 1, the range of the X values goes from -4 to 4.




Figure 2: CDF of Normal Distribution Curve

3.3 The Inverse Cumulative Distribution Function

Next, we use the inverse cumulative distribution function in the range of $[0,1)$ and a sample of 1000 times. We then move on to plot a histogram of 1000 of these observations. In order to find the inverse cumulative distribution, we use the **probability point function**, which returns the exact point where the probability of everything to the left is equal to the cumulative probability.

Lastly, we then move on to plotting the probability point function as a histogram. Using the `matplotlib` library, we then plot the probability value on the x-axis and the cumulative probability on the y-axis after applying the `.ppf`.

What we ended up getting is the graph as shown below in Figure 3. Here, the graph gives us the value of the random variable at a given cumulative probability. The x-axis represents represents the cumulative probability and the y-axis represents the frequency of the probabilities. We can observe that the graph gives the shape of a

normal distribution curve as a visible peak is seen when the probability is at 0.




Figure 3: PPF of Normal Distribution Curve

3.4 Normal Distribution Histogram of 10, 20, 40, 80 and 160 Samples

Here, we created a function that accepts 2 parameters, a `scipy` class distribution object and a integer `n`. Here, `n` is the number of samples to take from the distribution. This method should return a list of length `n` elements, where each element is the value of a single sample.

We use the function we created in the cell earlier to generate data and make histograms for each. Here are the details of the histograms as shown in the bullet lists and as Figures 4, 5, 6, 7 & 8 as shown in order.

- Normal Distribution of 10 Samples
- Normal Distribution of 20 Samples
- Normal Distribution of 40 Samples
- Normal Distribution of 80 Samples
- Normal Distribution of 160 Samples

3.4.1 Normal Distribution Histogram of 10 Samples

When we look at the normal distribution histogram of 10 samples, we can see that some of the values on the x-axis don't have any occurrences, as clearly displayed when observing the y-axis. Since this occurs at when the values are -0.25 and 0.25, we can describe the histogram as multi-modal due to 4 possible peak that occur at -0.75 , 0 , 0.35 and at 1 . It is quite difficult to deduce the skew of this histogram, as shown below as Figure 4.




Figure 4: Normal Distribution Histogram of 10 Samples

3.4.2 Normal Distribution Histogram of 20 Samples

When we look at the normal distribution histogram of 20 samples, we can see that there are a constant frequency between -2 to 2, excluding the downward peak at -1.5 to -1 and 1.5 to 2.0. Due to the upward peak, it can be safe to deduce that the histogram has a left skew, which is otherwise a negative skew, as shown in Figure 5.




Figure 5: Normal Distribution Histogram of 20 Samples

3.4.3 Normal Distribution Histogram of 40 Samples

When we look at the normal distribution histogram of 40 samples, we can see that there is a peak at the center at around 0. This peak means that the graph is unimodal and has no skew, as shown in Figure 6.

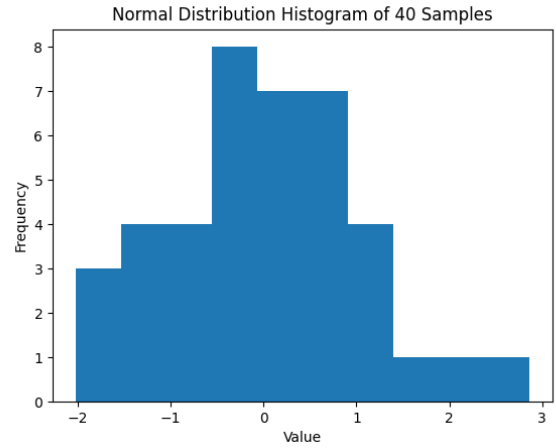


Figure 6: Normal Distribution Histogram of 40 Samples

3.4.4 Normal Distribution Histogram of 80 Samples

When we look at the normal distribution graph of 80 samples, we can see that there is a peak at -1 on the x-axis. We can see a second peak that occurs between 0 to 1. As a result, we can say that the histogram is bimodal. It is however, difficult to deduce the skew, as shown in Figure 7.




Figure 7: Normal Distribution Histogram of 80 Samples

3.4.5 Normal Distrubtion Histogram of 160 Samples

When we look at the normal distribution graph of 160 samples, we can see that there is a peak between 0 and 1 on the x-axis. As a result, we can deduce that this histogram is unimodal. When we take a look at the positioniing of the peak, it is observed to be more on the right side, and as result we can say that the histogram is left skewed, otherwise known as negatively skewed, as shown in Figure 8.




Figure 8: Normal Distribution Histogram of 160 Samples

3.5 Normal Probability Density Function over Histogram

For each one of the histograms, as shown in Figures 4 to 8, we directly plot the normal probability density function over each of the generated histograms. We achieve this by using the `seaborn` library. As a result, what we get is a Probability Density Function line being graphed as shown from Figures 9 to 13.

For each of thse figures, we can see that that the line being sketched is designed to fit as best as possible to the shape of the histogram, and in each of these figures, there is a visible peak that matches the peak of the histogram. Another common observation of the histograms from Figures 9 to 13 is that the x-axis ranges from around -3 to 3.

A notable exception to this is the one in Figure 9, which shows a peak over 0 on the x-axis, although there is no peak in the histogram.




Figure 9: Probability Density Function over Histogram of 10 Samples




Figure 10: Probability Density Function over Histogram of 20 Samples




Figure 11: Probability Density Function over Histogram of 40 Samples




Figure 12: Probability Density Function over Histogram of 80 Samples




Figure 13: Probability Density Function over Histogram of 160 Samples

3.6 Effect Size vs Iterations

We start off by drawing two sample groups using the same distribution and the same n , as follows:

- the first group should be 25 samples drawn from a normal distribution with mean = 0 and SD = 1
- the second group should be 25 different samples drawn from a normal distribution with mean = 0 and SD = 1

Next, we calculate the estimated effect size between the two groups. We calculate this by finding the absolute value of the difference between group 1 sample mean and group 2 sample mean.

We then created a function that accepts two distribution objects, two values of n for number of samples to draw from the distributions, and a integer k which is the number of times to loop the sample generations. This method should draw from each distribution n times and repeat this process for k iterations. As a result, this method returns three lists that are k elements long. The details of the return as follows:

- The elements in the first list are the estimated effect sizes at iteration k between the drawn samples of the two distributions

- Each element in the second list is a list of the drawn samples from distribution 1 at iteration k
- Each element in the third list is a list of the drawn samples from distribution 2 at iteration k

We then used our defined method to generate two sample groups, each with 25 observations for 1000 iterations. We then move to plot a histogram of the effect sizes as shown in Figure 14.

In this figure, we can see that the highest number of effect sizes range from 0 to 0.25, as they have the highest number of iterations. As a result, we get a peak, and the histogram is therefore right-skewed, which is otherwise positively skewed.

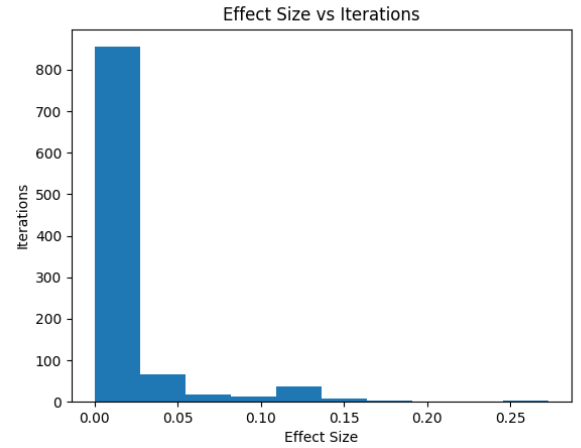


Figure 14: Effect Size vs Iterations

3.7 Histogram of Largest Effect Size of Sample Groups

Next, we find the index of the largest observed effect size. We use the index to plot the histograms of the two associated sample groups. As a result, we get Figure 15.

When we look at figure 15, we can see the color coded sample one and sampled, depicted as blue and orange. The blue signifies the first sample and the orange signifies the second sample. When we look at both of these histograms, we can deduce the peaks as follows:

- For Sample One, there are two peaks at 0 and 1, hence making the histogram bimodal
- For Sample Two, there are two peaks at -2 and 0, hence making the histogram bimodal




Figure 15: Histogram of Largest Effect Size of Two Sample Groups

3.8 Histogram of Effect Size

Next, we plot the histogram of the estimated effect sizes. We find the index of the minimum effect size and plot the two sample groups associated with this index using a histogram, as shown in Figure 16.

When we look at this histogram, we can deduce that the histogram is definitely unimodal and has a left skew, as a result making it negatively skewed.




Figure 16: Histogram of Effect Sizes

3.9 Histogram of Smallest Effect Size of Two Samples

We started off by instantiating the normal distributions. The first one should have the mean set to 0 and standard deviation set to 1. The second one should have the mean set to 1 and standard deviation set to 1. Next, we run 1000 iterations of sampling 25 values each.

This time, the first sample group should consist of 1000 iterations of 25 samples drawn from a normal distribution with $\mu = 0$ and $\sigma = 1$, and the second group should consist of 1000 iterations of 25 different samples drawn from a normal distribution with $\mu = 1$ and $\sigma = 1$.

As a result, we get Figure 17 as shown below. When we look at figure 17, we can see that there are two histograms, where the blue represents the Sample One and the orange represents the Sample Two. Here are a few observations:

- Sample One Histogram: it is a unimodal histogram that has no skew, the peak is at 0.5
- Sample Two Histogram: it is a unimodal histogram that has no skew, the peak is at 0.5 but at a Lower frequency, hence making it a lower peak.

4 Discussion

We now discuss our findings and conclude everything. We discuss with regards to our techniques and some of the observations that we have deduced with the data we produced and analyzed

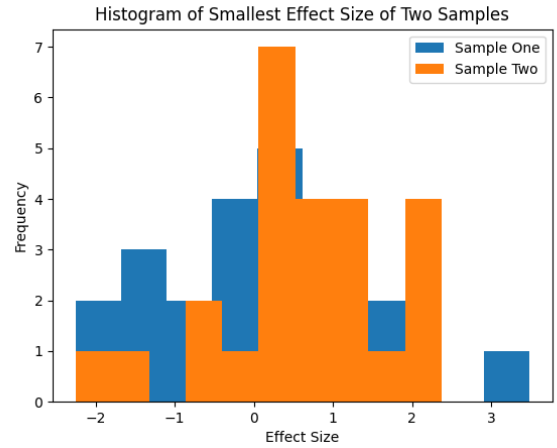


Figure 17: Histogram of Smallest Effect Size of Two Samples

4.1 Sampling and Normal Distribution

When we generated our sample values, we found the mean and normal distribution to be approximately zero, with a value of -4.547×10^{-16} . With the method we used, we believe that samples generated in this manner can approximate the true distribution well. In order to better job of approximating the true distribution, we can increase the size of the samples, as larger the sample size, the better the sample will approximate the true distribution.

4.2 The Sampling Method

Moving on, when we began plotting the normal distribution histograms, and as the samples increased from 10 to 160, we can see that the histograms begin to go from a bi-modal shape to unimodal shape, and as the sample increases from 20 to 160, the peak shifts from the right to the center.

4.3 The Type I and II Error

After we plotted the histograms, we wanted to look at type I and type II errors. They are explained in detail as such:

- Type I Error: it is when we reject a null hypothesis when it is true
- Type II Error: it when we don't reject the null hypothesis when it is false

After establishing an understanding of these errors, we drew two sample groups using the same normal distribution. The details of the two sample groups are:

- Sample Group One has 25 samples drawn from a normal distribution with mean 0 and standard deviation 1
- Sample Group Two has 25 different samples drawn from a normal distribution with mean 0 and standard deviation 1

We then calculated the effect size of those two sample groups, and these are the results, respectively:

- Mean of Random Sample One: 0.1057
- Mean of Random Sample Two: 0.1381
- Estimated Effect Size: 0.0324